



COMPUTATIONAL APPLICATIONS TO POLICY AND STRATEGY (CAPS)

Session 1 – Introduction to CAPS

Leo Klenner

Outline

1. Admin
2. About CAPS
3. Game Theory, Games and AI
4. StarCraft II
5. AI Policy and AI Strategy
6. Required Software



Workshop Information

- > Six sessions from 10/26 to TBA
- > Friday 4:00 to 5:00 pm
- > Course website: <https://git.io/fAaJn>
- > Contact:
 - > Leo (lklenne1@jhu.edu) # general
 - > Jonathan (jliu161@jhu.edu) # coding
 - > Henry (hfung4@jhu.edu) # data mining



Goals

- > Leverage state of the art research on AI gameplay to explore in open-ended manner how computational methods can advance IR research
- > Cover a lot of ground fast to provide you with a valuable short project and a personal website that augment your SAIS portfolio
- > Python is a tool not the primary object of learning itself
- > All of the sessions are prepared but we maximize value by working as a team



Goals – Class Project

- > Short technical report on any of the topics we discuss
- > Use code and software output or not
- > Examples: memo on constraints of deploying an AI system in a specific scenario, report on StarCraft II as an environment to conduct strategy research...
- > We will discuss the project in more detail after session 3



Core Resources

- > Readings and additional learning resources on our GitHub: <https://git.io/fAaJn>



Workshop Overview

- > Session 1: Introduction to CAPS and background on AI and games
- > Session 2: Primer on Python
- > Session 3: StarCraft II recap and building a rule-based bot
- > Session 4: Introduction to basic AI and building a learning-based bot
- > Session 5: Game data mining and discussion of AI Policy and Strategy
- > Session 6: Creating and hosting a free personal website through GitHub





About CAPS

- > How can computational methods advance our understanding and practice of international relations?
- > Why is it necessary that the IR community at large engages with emerging technologies such as AI?



The Intersection of AI and IR



What are the Consequences?

“The Age of Reason originated the thoughts and actions that shaped the contemporary world order. But that order is now in upheaval amid a new, even more sweeping technological revolution whose consequences we have failed to fully reckon with, and whose culmination may be a world relying on machines powered by data and algorithms and ungoverned by ethical or philosophical norms.”

- Henry Kissinger, *How the Enlightenment Ends* (June 2018)



How Companies React

“The only way through our “Crisis of Trust” is adopting a set of core values that allows us to navigate these complex times. We are now all stewards of the ethical & humane use of tech.”

- Mark Benioff, CEO Salesforce, Tweet (09/10/2018)

“I think one of the only ‘arms races’ AI people are excited about is the emerging competition between AI research labs to staff up meaningful policy organizations.”

- Jack Clark, Policy Lead OpenAI, Tweet (11/10/2018)



Who is in the AI/IR Market?

AI + AI Policy Labs



Venture and Partnerships



AI Policy Research




Oxford Digital Diplomacy
Research Group

External Stakeholders



Applications



Reassessing the Consequences

- > Do you think Kissinger's prediction is right?
- > What might be reasons to disagree with it?



Defining AI

> Russel and Norvig (1995)

- a. Systems that think like humans
- b. Systems that think rationally
- c. Systems that act like humans
- d. Systems that act rationally

> Obstacles

- a. How do humans solve problems?
- b. How to formalize uncertain knowledge?
- c. How to successfully pass the Turing Test?
- d. How to always do the right thing?



The Rational Agent Approach

> Properties

- a. Rational agency does not necessarily depend on correct inference
- b. Standard for rationality is clearly defined and completely general
- c. Limited rationality: can't always do the right thing in complicated environments


> Advantages

- a. Allows easy comparison between AI and non-AI agents
- b. Familiar from game theory and easily relatable to gameplay



Agents Interact with Environments

> General



> Prisoner's Dilemma

		P2	
		Cooperate	Defect
P1	Cooperate	(3,3)	(0,5)
	Defect	(5,0)	(1,1)



Prisoner's Dilemma – Agent

Properties of the PD agent	
Behavior	Maximize rewards
Lifetime	1 action
Personality	Homogenous
Memory	No memory
Strategies	Defector, Cooperator, Random, TFT, ...



Prisoner's Dilemma – Environment

Properties of the PD environment	
Type	Game theory, non-zero sum
Gameplay	Turntaking
Action space	2 = Cooperate, Defect
Environment states	4 = CC, CD, DC, DD
Rule of transition between states	Discrete, based on the agents' actions
Rewards assigned to each state	3/3, 5/0, 0/5, 1/1
Reward horizon	Instantaneous
Mode of information	Perfect Information
Mode of action	Simultaneous



An Interim Conclusion

- > Intelligent agents are conditioned by their environment
- > Consider the environment before you consider the agent
- > Games provide powerful, scalable environments




Games and AI

- > Two contemporary milestones in AI gameplay



- > 1997 Deep Blue beats Garry Kasparov
- > Algorithm: Alpha-Beta Search
- > Approach: Brute force



- > 2016 AlphaGo beats Lee Sedol
- > Algorithm: Hybrid (Neural nets, MCTS)
- > Approach: Deep learning



Video Games


- > Computationally more complex than Chess and Go
 - > Multi-objective tasks to reach goal
 - > Multiple dissimilar units
 - > Partial information
 - > Long time horizons
 - > Continuous action space and environment
- > Require teamplay, depending on the type of game
 - > Learning cooperation and teamplay is a paradigm shift for agents
 - > Opens a host of real-world applications




Recent AI Advances in Video Games



- > 2013 Atari DQN
- > DeepMind
- > Non-cooperative



- > 2018 Quake III Arena CTF
- > DeepMind
- > Hybrid



- > 2015 Malmo Minecraft
- > Microsoft
- > Cooperative



- > 2018 Five Dota
- > OpenAI
- > Hybrid

Enter StarCraft II



Screenshot from DeepMind's pyc2 API



Python for StarCraft II

- > 2017: release of the pyc2 and python-sc2 packages for StarCraft II
- > pyc2 (DeepMind)
 - > Vinyals, O., et al. 2017. *StarCraft II: A New Challenge For Reinforcement Learning*
 - > Geared towards building advanced reinforcement learning agents
- > python-sc2 (Dentosal)
 - > Geared towards ease of use for building both rule-based and AI agents



A Basic StarCraft II Agent

- > ‘Worker rush’
- > Simple example of a rule-based agent < 20 lines of code
- > Take everything you have at time t_0 and attack the enemy



- > At t_0 :
 - > Select all workers w_0
 - > Send w_0 to attack enemy start location

Code for Worker Rush

> Taken from <https://github.com/Dentosal/python-sc2#example>

```
> import sc2
> from sc2 import run_game, maps, Race, Difficulty
> from sc2.player import Bot, Computer

> class WorkerRushBot(sc2.BotAI):
    async def on_step(self, iteration):
        if iteration == 0:
            for worker in self.workers:
                await self.do(worker.attack(self.enemy_start_locations[0]))

> def main():
    run_game(maps.get("Simple128"), [
        Bot(Race.Terran, WorkerRushBot()),
        Computer(Race.Protoss, Difficulty.Easy)
    ], realtime = True)

> if __name__ == '__main__':
    main()
```



StarCraft II – Agent

Properties of the SCII agent	
Behavior	Win game
Lifetime	Until defeat
Personality	Asymmetric, 3 different races, multiple 100s of units
Memory	Depends
Strategies	Balance resource management, expanding vs. defense



StarCraft II – Environment


Properties of the SCII environment	
Type	Real-time strategy (RTS) game
Gameplay	Fast paced micro-actions and need for high-level planning
Action space*	10^8 , need for hierarchical actions
Environment states**	$10^{1,685}$
Rule of transition between states	Continuous, based on the agents' actions
Rewards assigned to each state	Unknown
Reward horizon	Long pay-off = strats more important than micro
Mode of information	Fog of war = imperfect information
Mode of action	Simultaneous

* Vinyals, O., et al. 2017. *StarCraft II: A New Challenge for Reinforcement Learning*. <https://arxiv.org/abs/1708.04782>

** Estimated for StarCraft Brood Wars. Usunier, N., et al. 2016. *Episodic Exploration for Deep Deterministic Policies: An Application to StarCraft Micromanagement Tasks*. <https://arxiv.org/abs/1609.02993>




Agent-Environment Loop 1



- > At each step t the agent
 - a. Executes action A_t
 - b. Receives observation O_t
 - c. Receives scalar reward R_t
- > The environment
 - a. Receives action A_t
 - b. Emits observation O_{t+1}
 - c. Emits scalar reward R_{t+1}
- > t increments at environment step



Agent-Environment Loop 2



- > At each step t the agent
 - a. Executes action A_t based on fixed policy
 - b. Receives observation O_t
- > The environment
 - a. Receives action A_t
 - b. Emits observation O_{t+1}
- > t increments at environment step
- > Observations restricted by fixed policy



Differences between Loop 1 and Loop 2

> Loop 1

- a. Reward enables utility max.
- b. Utility max. enables learning
- c. Learning enables dynamic actions
- d. Dynamic actions enable adaption
- e. Reinforcement learning

> Loop 2

- a. No reward, hence no utility
- b. No utility, hence no learning
- c. No learning, hence static actions
- d. Static actions, hence no adaption
- e. Rule-based



Respective Pros and Cons

> Loop 1

- a. Accommodates complex actions
- b. Requires meaningful rewards
- c. Underperforms in 'mute' env.
- d. Weakly predictable actions

> Loop 2

- a. Complexity constrained by coded rules
- b. Does not depend on rewards
- c. Constant actions across environments
- d. 100% predictable actions

> In theory, learning > rule-based. In practice, heavily dependent on environment.



Hard Trade-Offs

- > Assuming that you have only limited knowledge about the environment and only one shot at success, which approach would you chose? Why?
- > Would your answer change if you had more than one shot at success?



Framing the Trade-Offs

- > The global implementation of AI brings about a host of questions
- > AI Policy and AI Strategy have emerged as fields to provide answers
- > **AI Policy**
 - > Analysis and practice of societal decision-making about AI
- > **AI Strategy**
 - > Long-term conceptual analysis of AI developments
- > While AI researchers recognize IR as a valuable discipline, the IR community has yet to leverage its potential.



Challenge

- > We need to integrate both the computational tools and the debate on implementing AI into IR thinking
- > What are useful points of contact between AI and IR?
- > Where do you see overlap and potential for applications?



Summary

- > AI is having an increasing impact on the global order
- > Stakeholders are looking for meaningful ways to respond
- > Clear definitions of AI are needed to mitigate uncertainty
- > The rational actor approach to AI offers high explanatory power
- > Games and game theory can be used to develop AI agents
- > Current video games present highly complex challenges for AI
- > One of the games with the most strategic depth is StarCraft II
- > Two Python APIs allow building AI and rule-based agents for StarCraft II
- > The IR community needs to integrate these tools into its thinking and practice



Required Software

- > For next week you only need Python. We will install more stuff along the way.
- > Python 3.6.6
 - > <https://www.python.org/downloads/release/python-366/>
- > Python-sc2
 - > via command line: pip install sc2
- > StarCraft II (free to play)
 - > <https://us.battle.net/account/download/>
 - > This will take some time so best to do it at home
- > StarCraft II maps
 - > <https://github.com/Blizzard/s2client-proto#map-packs>
 - > Once installed, extract the maps as subdirectories into StarCraft II's map directory
 - > Do this after you have installed StarCraft II

